

# **Tor Bridges: Hiding Public Resources from Nation-State Adversaries**

Isis Agora Lovecraft  
The Tor Project

**What is Tor?**

# How Tor Works

# Tor Bridges

---

- Secret entrances to the Tor network
- Must be distributed out-of-band
- DPI or an active adversary is required to identify Bridges
- Distributed via a centralised Bridge distribution system called BridgeDB

# Arms Race

# Arms Race

---

- Since 2010: China's GFW began active probing Tor Bridges
  - Observe Tor client's TCP connection to the Bridge
  - For Tor<0.2.3.17-beta, identification was based upon Tor's unique TLS ciphersuite list
  - A seemingly random machine from somewhere in China (possibly using IP-spoofing) will connect to the Bridge's IP:port and attempt to complete the first couple steps of the handshake
  - The Bridge is usually blocked by IP:port
  - The GFW sometimes spoofs a RST from Bridge to the client

# Arms Race

---

- 2012: Ethiopia began blocking all TLS by looking for the client HELLO.

Any packet with the string

**TLS\_DHE\_RSA\_WITH\_AES\_256\_CBC\_SHA**

in it is dropped...

...but when you specify a preference for

**TLS\_DHE\_RSA\_WITH\_AES\_128\_CBC\_SHA**

instead (or split the ciphersuite list) it works just fine.

# Arms Race

```
// START_DEFINITION
requires grammar version 5
/**
 * Identify clients accessing Tor bridge information.
 */
fingerprint('anonymizer/tor/bridge/tls') =
ssl_x509_subject('bridges.torproject.org') or
ssl_dns_name('bridges.torproject.org');

/**
 * Database Tor bridge information extracted from confirmation emails.
 */
fingerprint('anonymizer/tor/bridge/email') =
email_address('bridges@torproject.org')
and email_body('https://bridges.torproject.org/' : c++
extractors: {{
    bridges[] = /bridge\s([0-9]{1,3}\. [0-9]{1,3}\. [0-9]{1,3}\. [0-9]{1,3})?:?([0-9]{2,4}?[^\0-9])?/;
}}
init: {{
    xks::undefine_name("anonymizer/tor/torbridges/emailconfirmation");
}}
main: {{
    static const std::string SCHEMA_OLD = "tor_bridges";
    static const std::string SCHEMA_NEW = "tor_routers";
    static const std::string FLAGS = "Bridge";
    if (bridges) {
        for (size_t i=0; i < bridges.size(); ++i) {
            std::string address = bridges[i][0] + ":" + bridges[i][1];
            DB[SCHEMA_OLD]["tor_bridge"] = address;
            DB.apply();
            DB[SCHEMA_NEW]["tor_ip"] = bridges[i][0];
            DB[SCHEMA_NEW]["tor_port_or"] = bridges[i][1];
            DB[SCHEMA_NEW]["tor_flags"] = FLAGS;
            DB.apply();
        }
        xks::fire_fingerprint("anonymizer/tor/directory/bridge");
    }
    return true;
}});
// END_DEFINITION
```

2014: Das Erste releases several NSA XKEYSCORE selectors. These are the XKEYSCORE selectors for Bridge IPs handed out from the current bridge distribution systems.

<http://daserste.ndr.de/panorama/xkeyscorerules100.txt>



**Pluggable Transports,  
obfsproxy, & other  
completely  
unpronounceable things**

# obfs4proxy

---

- Tor's NTor handshake with public keys obfuscated via the Elligator 2 mapping
- Link layer uses NaCl secret boxes (Poly1305, XSalsa20)

# Simple Formulae

---

- Make the handshake as uniform as possible
- Use some pre-shared key material for authentication of the server and encrypt starting with the client's first message

**This just sweeps the  
problems under the rug**

# Bridge Distribution

---

Bridge obfs4 106.187.37.158:62421 50182425F17DEF0B51B0790188D2E04E300314B7  
cert=pKDDKPfTYDJjX2tJbm6z/CW3+dnEg1vw3YjofAw2fbDnHJ2Rc7/yTAFg/1RiyoMme5Dgcw iat-  
mode=0

Bridge obfs4 178.209.52.110:443 67E72FF33D7D41BF11C569646A0A7B4B188340DF  
cert=Z+cv8z19Qb8RxWlkagp7SxiDQN++b7D2Tntowhf+j4D15/kLuj3EoSSGvuREGPc3h600fw iat-  
mode=0

Bridge obfs4 83.212.101.3:41213 A09D536DD1752D542E1FBB3C9CE4449D51298239  
cert=lPRQ/MXdD1t5SRZ9MquYQNT9m5DV757jtdXdlePmRCudUU9CFU0X1Tm7/meFSyP0sud7Cw iat-  
mode=0

We need some manner of getting these  
Bridge connection details and pre-shared  
key material to Tor clients, and only to Tor  
clients.

# Proof-of-Work Schemes Won't Work

---

- The adversary is always going to have access to significantly more resources than the Tor users we're trying to help.

This applies to anything:

- Resource-based (CPU and/or memory)
- Monetary (*Yes, even Bitcoin...*)
- CAPTCHA-solving capabilities
- The thirty other PoW schemes proposed to me over the years.

**If the adversary is so  
powerful, then what *can't*  
they do?**

**If the adversary is so  
powerful, then what *can't*  
they do?**

**Make friends!**



**What if we had a system  
where well-behaved users  
could invite their friends to  
get Bridges?**

With this, we solve one problem and introduce another: the Bridge distribution system now has the complete social graph of which Tor Bridge users know each other, making it an even more attractive target for nation-state adversaries.


# rBridge

---


Wang, Q., Lin, Z., Borisov, N., & Hopper, N. (2013, February). rBridge: User Reputation based Tor Bridge Distribution with Privacy Preservation. In *NDSS*.

- Users are given “brownie points” for “good behaviour”
- Users with enough brownie points might win the chance to invite their friends
- Censors lock themselves out of the system via their own bad behaviour. Also nobody wants to be friends with those losers anyway.


# The Best Strategy for Censors



(a) Unblocked bridges



(b) Thirsty users



(c) Thirsty users with backup bridges

Figure 4: Event-driven blocking ( $f = 5\%$ )

# Some Minor Problems

---

Some odd crypto choices, silly mistakes, and efficiency sacrifices for very little added privacy

- K-TAA signature scheme O\_o'
- Pedersen commitments on vectors
- Oblivious Transfer
- Ad-hoc anonymous credential construction from k-TAA signatures and Camenisch-Stadler NIZK proof of discrete logarithm.

# Redesign

---

- CL Anonymous Credentials

Belenkiy, M., Camenisch, J., Chase, M., Kohlweiss, M., Lysyanskaya, A., & Shacham, H. (2009). Randomizable proofs and delegatable anonymous credentials. In Advances in Cryptology-CRYPTO 2009 (pp. 108-125). Springer Berlin Heidelberg.

Uses Boneh-Boyen signatures and rerandomised Groth-Sahai NIZK proofs of satisfiability of a pairing-product equation to construct delegatable authentication tokens with attributes.

# Redesign

---

- CL Anonymous Credentials

Removes the need for k-TAA signatures, fixes the issues with making Pedersen commitments on vectors of independent variables, and provides a slightly less insane alternative to the credential constructions.

# Redesign

$$\pi_2 = \text{NIPK} \left\{ \begin{array}{l} (x, \Phi, C_\Phi, e_\Phi, s_\Phi, s'_\Phi, r_\Phi^{(1)}, r_\Phi^{(2)}, \delta_\Phi^{(1)}, \delta_\Phi^{(2)}, \\ B_u, \tau_u, \phi_u, C_u, e_u, s_u, s'_u, r_u^{(1)}, r_u^{(2)}, \delta_u^{(1)}, \delta_u^{(2)}, \\ \bar{\Phi}, \bar{s}'_\Phi, \bar{\phi}_u, \bar{s}'_u) : \\ \bigwedge_{j=1}^m [b_j \neq z^{B_u}] \wedge \\ C_u = g_1^{s'_u} g_2^x g_3^{B_u} g_4^{\tau_u} g_5^{\phi_u} \wedge \\ A_u^{(1)} = g_1^{r_u^{(1)}} g_2^{r_u^{(2)}} \wedge \\ (A_u^{(1)})^{c_u} = g_1^{\delta_u^{(1)}} g_2^{\delta_u^{(2)}} \wedge \\ \frac{\hat{c}(A_u^{(2)}, pk)}{\hat{c}(g_0, h)} = \hat{e}(A_u^{(2)}, h)^{-c_u} \hat{e}(g_2, y)^{r_u^{(1)}} \\ \hat{e}(g_2, h)^{\delta_u^{(1)}} \hat{e}(g_1, h)^{s_u} \hat{e}(g_2, h)^x \\ \hat{e}(g_3, h)^{B_u} \hat{e}(g_4, h)^{\tau_u} \hat{e}(g_5, h)^{\phi_u} \wedge \\ C_\Phi = g_1^{s'_\Phi} g_2^x g_3^\Phi \wedge \\ A_\Phi^{(1)} = g_1^{r_\Phi^{(1)}} g_2^{r_\Phi^{(2)}} \wedge \\ (A_\Phi^{(1)})^{c_\Phi} = g_1^{\delta_\Phi^{(1)}} g_2^{\delta_\Phi^{(2)}} \wedge \\ \frac{\hat{c}(A_\Phi^{(2)}, pk)}{\hat{c}(g_0, h)} = \hat{e}(A_\Phi^{(2)}, h)^{-c_\Phi} \hat{e}(g_2, y)^{r_\Phi^{(1)}} \\ \hat{e}(g_2, h)^{\delta_\Phi^{(1)}} \hat{e}(g_1, h)^{s_\Phi} \hat{e}(g_2, h)^x \hat{e}(g_3, h)^\Phi \wedge \\ \kappa_\Phi = z^{s_\Phi} \wedge \\ t_u = T_{cur} - \tau_u \wedge \\ \left[ (t_u < T_0 \wedge \bar{\phi}_u = 0) \vee \right. \\ \left. (t_u \geq T_0 \wedge t_u \leq T_1 \wedge \bar{\phi}_u = \rho(t - T_0)) \vee \right. \\ \left. (t_u > T_1 \wedge \bar{\phi}_u = \rho(T_1 - T_0)) \right] \wedge \\ \bar{\Phi} = \Phi + \bar{\phi}_u - \phi_u \wedge \\ \bar{C}_u = g_1^{\bar{s}'_u} g_2^x g_3^{B_u} g_4^{\tau_u} g_5^{\bar{\phi}_u} \wedge \\ \bar{C}_\Phi = g_1^{\bar{s}'_\Phi} g_2^x g_3^\Phi \wedge \end{array} \right\}$$

$$\pi_3 = \text{NIPK} \left\{ \begin{array}{l} (x, \Phi, C_\Phi, e_\Phi, s_\Phi, s'_\Phi, r_\Phi^{(1)}, r_\Phi^{(2)}, \delta_\Phi^{(1)}, \delta_\Phi^{(2)}, \tau_b, \\ \phi_b, C_b, e_b, s_b, s'_b, r_b^{(1)}, r_b^{(2)}, \delta_b^{(1)}, \delta_b^{(2)}, \bar{\Phi}, \bar{s}'_\Phi) : \\ C_b = g_1^{s'_b} g_2^x g_3^{B_b} g_4^{\tau_b} g_5^{\phi_b} \wedge \\ A_b^{(1)} = g_1^{r_b^{(1)}} g_2^{r_b^{(2)}} \wedge \\ (A_b^{(1)})^{c_b} = g_1^{\delta_b^{(1)}} g_2^{\delta_b^{(2)}} \wedge \\ \frac{\hat{c}(A_b^{(2)}, pk)}{\hat{c}(g_0, h)} = \hat{e}(A_b^{(2)}, h)^{-c_b} \hat{e}(g_2, y)^{r_b^{(1)}} \\ \hat{e}(g_2, h)^{\delta_b^{(1)}} \hat{e}(g_1, h)^{s_b} \hat{e}(g_2, h)^x \\ \hat{e}(g_3, h)^{B_b} \hat{e}(g_4, h)^{\tau_b} \hat{e}(g_5, h)^{\phi_b} \wedge \\ \kappa_b = z^{s_b} \wedge \\ C_\Phi = g_1^{s'_\Phi} g_2^x g_3^\Phi \wedge \\ A_\Phi^{(1)} = g_1^{r_\Phi^{(1)}} g_2^{r_\Phi^{(2)}} \wedge \\ (A_\Phi^{(1)})^{c_\Phi} = g_1^{\delta_\Phi^{(1)}} g_2^{\delta_\Phi^{(2)}} \wedge \\ \frac{\hat{c}(A_\Phi^{(2)}, pk)}{\hat{c}(g_0, h)} = \hat{e}(A_\Phi^{(2)}, h)^{-c_\Phi} \hat{e}(g_2, y)^{r_\Phi^{(1)}} \\ \hat{e}(g_2, h)^{\delta_\Phi^{(1)}} \hat{e}(g_1, h)^{s_\Phi} \hat{e}(g_2, h)^x \hat{e}(g_3, h)^\Phi \wedge \\ \kappa_\Phi = z^{s_\Phi} \wedge \\ t_b = \beta_b - \tau_b \wedge \\ \left[ (t_b < T_0 \wedge \bar{\phi}_b = 0) \vee \right. \\ \left. (t_b \geq T_0 \wedge t_b \leq T_1 \wedge \bar{\phi}_b = \rho(t_b - T_0)) \vee \right. \\ \left. (t_b > T_1 \wedge \bar{\phi}_b = \rho(T_1 - T_0)) \right] \wedge \\ \bar{\Phi} = \Phi + \bar{\phi}_b - \phi_b - \phi^- \wedge \\ \bar{\Phi} > 0 \\ \bar{C}_\Phi = g_1^{\bar{s}'_\Phi} g_2^x g_3^\Phi \wedge \end{array} \right\}$$

# Remove Oblivious Transfer

---

- rBridge uses  $n$ -out-of- $m$  OT to hide which Bridges are distributed to a client at the time of distribution.
- Another messy construction for an additional proof of inequality of openings to commitments to chosen Bridge and some previous Bridge to avoid duplicates.
- In the end, the client tells the server which Bridge it has when it reports the Bridge was blocked.



# Open Questions

---

- What do we mean when we say “A Bridge is blocked”? Does “blocked in China” mean “blocked in Iran”? What if China sells data on Tor Bridges to Iran?
- Simpler anonymous credential constructions which don't require pairings? E.g. based upon algebraic MACs or Diffie-Hellman.

**All these improvements  
only protect the Bridges  
from an adversary on the  
client side.**

**What other ways are  
there to discover Tor  
Bridges?**

# Arms Race

---

- 2015: Some researchers in China realise there's another way...

## Extensive Analysis and Large-Scale Empirical Evaluation of Tor Bridge Discovery

Zhen Ling\*, Xinwen Fu<sup>†</sup>, Wei Yu<sup>‡</sup>, Junzhou Luo\* and Ming Yang\*

\*Southeast University, Email: {zhenling,jluo,yangming2002}@seu.edu.cn

<sup>†</sup>University of Massachusetts Lowell, Email: xinwenfu@cs.uml.edu

<sup>‡</sup>Towson University, Email: wyu@towson.edu

# Arms Race

---

and at an apartment. We formally analyzed catch probability of discovering bridges via middle onion routers. Our real-world implementation and large-scale experiments validated effectiveness and feasibility of the three bridge discovery approaches. **We have discovered 2365 Tor bridges via email and https and 2369 bridges by only one controlled Tor middle router in 14 days.** The malicious middle router based bridge discovery approach is simple, incurs little overhead, can discover bridges distributed by any approach, and is efficient and effective. We also discussed potential mechanisms to

# Arms Race

---



# Solution

---

- Tor Proposal #188:

<https://gitweb.torproject.org/torspec.git/tree/proposals/188-bridge-guards.txt>

- A Tor client chooses three hops through the Tor network.
- The client's traffic goes through those three hops. Well... yes. But...
- Those aren't the only hops which the client's traffic can go through.

# Solution

---

- Nothing in Tor's circuit-level crypto prevents those three hops from adding more hops to the client's path. Tor turns out to be loose-source routed.
- The client doesn't even have to know it's happening.
- We can use this to have Tor Bridges choose a Bridge Guard node, and transparently inject the Bridge's Guard into the client's path.



# Questions